

COMP0197 CW2

This codebase contains submission code for COMP0197 Applied Deep Learning coursework 2.

Setting up Environment

Prerequisites

- **Python:** Version 3.x recommended.
- **pip:** Python package installer.
- **Git:** For cloning the repository (if not already downloaded).
- **(Optional):** CUDA-enabled GPU for significantly faster training and evaluation.

Base Environment

To create a base conda environment, please execute the below script.

```
conda create -n comp0197-cw2-pt python=3.12 pip && conda activate comp0197-cw2-pt && pip install torch==2.5.0 torchvision --index-url https://download.pytorch.org/whl/cpu
```

Extra Dependencies

In addition to above base packages, this codebase also requires below packages to be installed.

```
pip install scikit-image torchmetrics
```

Script Execution Instructions

Fully Supervised Model

Fully supervised models are trained to provide comparison with weakly supervised models. To train a fully supervised segmentation model:

```
python -m supervised.train --model=segnet
```

If you're using a MacBook with an M1, M2, or M3 chip, we suggest enabling CPU fallback for better compatibility:

```
PYTORCH_ENABLE_MPS_FALLBACK=1 python -m supervised.train --model=segnet
```

Available segmentation model options are: segnet, segnext, effunet, unet.

Weakly Supervised Model

NOTE: Below command are designed to be executed in sequence. Breaking the sequence may result in receiving errors due to saved model weights not found.

1. Finetune classifier without regularisation

```
python -m cam.finetune
```

or with regularisation

```
python -m crm.train_cam_with_crm
python -m crm.evaluate_with_crm
```

- If Mac users ran into segmentation fault, please download saved models on OneDrive
 - Resnet:

<https://1drv.ms/u/c/2ef0e412637ecc3c/EawGxav3g3BPke8uXA7C5W0Bdf2oIHQSoV6smZgRWXR1NA>
 - Reconstruction network:

https://1drv.ms/u/c/2ef0e412637ecc3c/EdhrCblkW6dEpXflmbAcRsoBBb_3ceJHz16NxfTiqLPmhg?e=llvOmN
 - Resnet DRS: <https://1drv.ms/u/c/2ef0e412637ecc3c/EQU-6ec3hklKHi9hTXwXxDEBWx5czmOywwLiH3gsT0qhAQ?e=nkVYyUQ>.
 - Reconstruction network Drs:

<https://1drv.ms/u/c/2ef0e412637ecc3c/EesRuHMqxxgZAvj6Qc710poYBfyskimMUQtJAFrC9wmOCw?e=YEFLuq>
 - Please place them under `./crm_models/`

Note: `crm.train_cam_with_crm` will generate superpixels under directory `./superpixels`. To skip the generation, simply download from OneDrive and unzip it under project root.

- Superpixels:

<https://1drv.ms/u/c/2ef0e412637ecc3c/EQy9SXX7x4tGnqJWRplJa7EBYK9I7c2ipQB07oCzcjAfKQ?e=Ovlgq5>
- If training was killed early, rerunning could lead to into `superpixels/xxx not found error`, please try to remove the `./superpixels` folder and rerun, or download and zip from OneDrive.

All available command argument for `train_cam_with_crm`:

- `--model`: Classification model name, allowed value ['resnet', 'resnet_drs']
- `--cls_lr`: Classifier learning rate

- `--rec_lr`: Reconstruction net learning rate
- `--vgg_weight`: Weight for VGG loss
- `--align_weight`: Weight for alignment loss
- `--epochs`: Number of training iteration

All available command argument for `evaluate_with_crm`:

- `--model`: Classification model name, allowed value ['resnet', 'resnet_drs']

2. Get pseudo masks generated by extracting CAM from the classifier:

```
python -m cam.postprocessing --model=resnet
```

or with regularisation

```
python -m cam.postprocessing --model=resnet_crm
```

All available command argument for `postprocessing`:

- `--model`: Classification model name, allowed value ['resnet', 'resnet_crm', 'resnet_drs']

Sample heatmap images are generated in the `cam/output/cam_grid.jpg`

3. Train the SegNet segmentation model with pseudo masks

If classifier used in previous steps is simple resnet, model and pseudo masks are save in `cam/saved_models` folder:

```
python -m supervised.train --model=segnet --pseudo --  
pseudo_path=cam/saved_models/resnet_pet_cam_pseudo.pt
```

If classifier used earlier is with regularisation, i.e. `resnet_crm` or `resnet_drs`, model and pseudo masks are save in `crm_models/` folder:

```
python -m supervised.train --model=segnet --pseudo --  
pseudo_path=cam/saved_models/resnet_pet_gradcampp_crm_pseudo.pt
```

All available command argument for `train`:

- `--model`: Segmentation model name, allowed value ['segnet', 'segnext', 'effunet', 'unet']
- `--pseudo`: boolean, whether to use pseudo mask
- `--pseudo_path`: path where pseudo mask is saved
- `--verbose`: boolean, whether to print verbose message

Ablation Experiments

Ensure the classification model is saved under `cam/saved_models`

```
python -m ablation.grid_search
```

All available command argument for `grid_search`:

- `--model`: Segmentation model name, allowed value ['segnet', 'segnext', 'effunet', 'unet']
- `--model_path`: Path where trained classifier model is saved
- `--result_path`: Path to search results.

Open-ended Section

This section details the steps to download data, prepare labels (if needed), train models with different supervision strategies, and evaluate them.

Step 1: Download Data

This script checks if the Oxford-IIIT Pet Dataset has been downloaded already. If not, download it.

```
python -m open_ended.download_data
```

- This command will download the dataset into a directory (likely `./data` based on subsequent commands). Ensure you have sufficient disk space and an internet connection.

Step 2: Generate Weak Labels (Optional, Attached in submission)

This script generates the sparse weak annotations (points, scribbles, boxes) from the ground-truth segmentation masks provided by the original dataset.

```
python -m open_ended.weak_label_generator --data_dir ./data/oxford-iiit-pet --  
output_file ./open_ended/weak_labels/weak_labels_train.pkl
```

- `--data_dir ./data/oxford-iiit-pet`: Specifies the directory where the Oxford-IIIT Pet dataset was downloaded (input).
- `--output_file ./open_ended/weak_labels/weak_labels_train.pkl`: Specifies the path to save the generated weak labels as a Python pickle file (output).

Important Note: As mentioned in the original README, pre-generated weak labels (`weak_labels_train.pkl`) are already included in the `./open_ended/weak_labels/` directory within the repository. **Therefore, running this step is generally NOT necessary unless you want to regenerate the labels with different parameters or settings.**

Step 3: Train Segmentation Models

This is the core step where the segmentation model is trained using different weak supervision configurations. The script `open_ended/train.py` is used repeatedly with different arguments.

Common Training Arguments:

- `--supervision_mode [mode]`: **Crucial argument.** Specifies the type of weak supervision to use. Examples below use `points`, `scribbles`, `boxes`, and various `hybrid_...` combinations.
- `--run_name [name]`: A unique name for this specific training run. Used for logging and naming checkpoint files (e.g., `segnet_points_run1`).
- `--data_dir ./data/oxford-iiit-pet`: Path to the dataset directory.
- `--weak_label_path ./open_ended/weak_labels/weak_labels_train.pkl`: Path to the file containing the weak labels.
- `--batch_size 64`: Number of samples per batch during training. Adjust based on GPU memory.
- `--lr 2e-4`: Learning rate for the optimizer.
- `--epochs 25`: Number of training epochs.
- `--num_workers 8`: Number of worker processes for data loading. Adjust based on your system's CPU cores.
- `--img_size 256`: Resize input images to this square dimension.
- `--checkpoint_dir [path]`: Directory where model checkpoints (saved model weights) will be stored.

Training Examples:

(a) Training with Single Weak Supervision Types: These commands train separate models, each using only one type of weak annotation. Checkpoints are saved in `./checkpoints_single`.

- **Using Points:**

```
python -m open_ended.train \
  --supervision_mode points \
  --run_name segnet_points_run1 \
  --data_dir ./data/oxford-iiit-pet \
  --weak_label_path ./open_ended/weak_labels/weak_labels_train.pkl \
  --batch_size 64 \
  --lr 2e-4 \
  --epochs 25 \
  --num_workers 8 \
  --img_size 256 \
  --checkpoint_dir ./checkpoints_single
```

- **Using Scribbles:**

```
python -m open_ended.train \
  --supervision_mode scribbles \
  --run_name segnet_scribbles_run1 \
  --data_dir ./data/oxford-iiit-pet \
  --weak_label_path ./open_ended/weak_labels/weak_labels_train.pkl \
  --batch_size 64 \
  --lr 2e-4 \
```

```
--epochs 25 \  
--num_workers 8 \  
--img_size 256 \  
--checkpoint_dir ./checkpoints_single
```

- **Using Bounding Boxes:**

```
python -m open_ended.train \  
--supervision_mode boxes \  
--run_name segnet_boxes_run1 \  
--data_dir ./data/oxford-iiit-pet \  
--weak_label_path ./open_ended/weak_labels/weak_labels_train.pkl \  
--batch_size 64 \  
--lr 2e-4 \  
--epochs 25 \  
--num_workers 8 \  
--img_size 256 \  
--checkpoint_dir ./checkpoints_single
```

(b) Training with Hybrid Weak Supervision Types: These commands train models using combinations of weak annotation types. Checkpoints are saved in `./checkpoints_hybrid`.

- **Using Points + Scribbles:**

```
python -m open_ended.train \  
--supervision_mode hybrid_points_scribbles \  
--run_name segnet_hybrid_points_scribbles_run1 \  
--data_dir ./data/oxford-iiit-pet \  
--weak_label_path ./open_ended/weak_labels/weak_labels_train.pkl \  
--batch_size 64 \  
--lr 2e-4 \  
--epochs 25 \  
--num_workers 8 \  
--img_size 256 \  
--checkpoint_dir ./checkpoints_hybrid
```

- **Using Points + Boxes:**

```
python -m open_ended.train \  
--supervision_mode hybrid_points_boxes \  
--run_name segnet_hybrid_points_boxes_run1 \  
--data_dir ./data/oxford-iiit-pet \  
--weak_label_path ./open_ended/weak_labels/weak_labels_train.pkl \  
--batch_size 64 \  
--lr 2e-4 \  
--epochs 25 \  
--num_workers 8 \  
--img_size 256 \  
--checkpoint_dir ./checkpoints_hybrid
```

- **Using Scribbles + Boxes:**

```
python -m open_ended.train \  
  --supervision_mode hybrid_scribbles_boxes \  
  --run_name segnet_hybrid_scribbles_boxes_run1 \  
  --data_dir ./data/oxford-iiit-pet \  
  --weak_label_path ./open_ended/weak_labels/weak_labels_train.pkl \  
  --batch_size 64 \  
  --lr 2e-4 \  
  --epochs 25 \  
  --num_workers 8 \  
  --img_size 256 \  
  --checkpoint_dir ./checkpoints_hybrid
```

- **Using Points + Scribbles + Boxes:**

```
python -m open_ended.train \  
  --supervision_mode hybrid_points_scribbles_boxes \  
  --run_name segnet_hybrid_points_scribbles_boxes_run1 \  
  --data_dir ./data/oxford-iiit-pet \  
  --weak_label_path ./open_ended/weak_labels/weak_labels_train.pkl \  
  --batch_size 64 \  
  --lr 2e-4 \  
  --epochs 25 \  
  --num_workers 8 \  
  --img_size 256 \  
  --checkpoint_dir ./checkpoints_hybrid
```

Step 4: Evaluate Trained Models

After training, use the [open_ended/evaluate.py](#) script to evaluate the performance of the saved model checkpoints on the test set.

```
python -m open_ended.evaluate \  
  --data_dir ./data/oxford-iiit-pet \  
  --model_paths checkpoints_single/segnet_boxes_run1_best_acc.pth \  
               checkpoints_single/segnet_points_run1_best_acc.pth \  
               checkpoints_single/segnet_scribbles_run1_best_acc.pth \  
  
checkpoints_hybrid/segnet_hybrid_points_boxes_run1_best_acc.pth \  
  
checkpoints_hybrid/segnet_hybrid_points_scribbles_boxes_run1_best_acc.pth \  
  
checkpoints_hybrid/segnet_hybrid_points_scribbles_run1_best_acc.pth \  
  
checkpoints_hybrid/segnet_hybrid_scribbles_boxes_run1_best_acc.pth \  
  --batch_size 8 \  
  --device cuda
```

- `--data_dir ./data/oxford-iiit-pet`: Path to the dataset directory.
- `--model_paths [path1] [path2] ...`: **Important:** List the paths to the specific checkpoint files (`.pth`) you want to evaluate. These should typically be the checkpoints saved based on the best validation performance during training (e.g., `_best_acc.pth` or similar, as indicated by the filenames). Ensure these paths correctly point to the files generated in Step 4.
- `--batch_size 8`: Batch size for evaluation. Can often be larger than training batch size depending on GPU memory.
- `--device cuda`: Specifies the device for evaluation. Use `cuda` for GPU or `cpu` for CPU.

(Note: The second `evaluate.py` command listed in the original README under a separate "Evaluate" heading seems potentially inconsistent or refers to different models/runs not shown in the main training sequence. This rewritten guide focuses on evaluating the models trained in Step 4 above.)
